

4.1. Create a Custom IAM Policy

Created IAM policy which grants list and get permissions to lab bucket

As seen In AWS Console:

Policy name	Type	Used as	Description
s3-list-get-discordlabs-ec2	Customer managed	Permissions policy...	grants List and Get permissions to the discord-labs bucket

Policy:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": "s3:ListBucket",
      "Resource": [
        "arn:aws:s3:::[REDACTED]"
      ]
    },
    {
      "Effect": "Allow",
      "Action": "s3:GetObject",
      "Resource": [
        "arn:aws:s3:::[REDACTED]/*"
      ]
    }
  ]
}
```

4.2. Create an IAM Role for EC2


Created IAM role and attached policy to it


EC2-S3-Access-Role-Brandon [Info](#) [Delete](#)


Allows EC2 instances to call AWS services on your behalf.

Summary [Edit](#)

Creation date
June 13, 2025, 10:42 (UTC-07:00)

ARN


Instance profile ARN
 profile/
EC2-S3-Access-role-Brandon

Last activity
 20 days ago


Maximum session duration
1 hour


[Permissions](#) [Trust relationships](#) [Tags](#) [Last Accessed](#) [Revoke sessions](#)

Permissions policies (1) [Info](#) [Refresh](#) [Simulate](#) [Remove](#) [Add permissions](#)

You can attach up to 10 managed policies.

Filter by Type
All types

< 1 > 

<input type="checkbox"/>	Policy name ?	Type	Attached entities
<input type="checkbox"/>	 s3-list-get-discordlabs-ec2	Customer managed	1

4.3. Attach the Role to Your EC2 Instance

I previously attached a role to the EC2 in lab2, and needed to replace it with the newly created IAM role above. First, had to get the existing AssociationId (using describe-iam-instance-profile-associations)

```
aws ec2 describe-iam-instance-profile-associations \
--filters Name=instance-id,Values=i-0d0a94a78842272ff
```

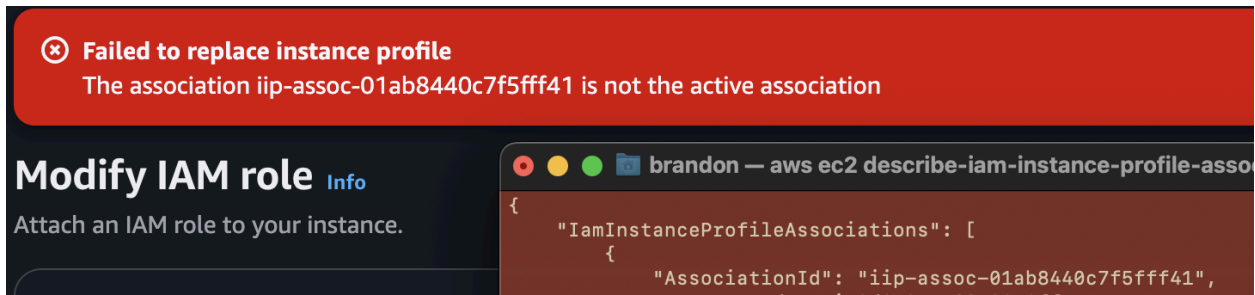
```
{
  "IamInstanceProfileAssociations": [
    {
      "AssociationId": "iip-assoc-01ab8440c7f5fff41",
      "InstanceId": "",
      "IamInstanceProfile": {
        "Arn": "",
        "Id": ""
      },
      "State": "associated"
    }
  ]
}
```

(END)

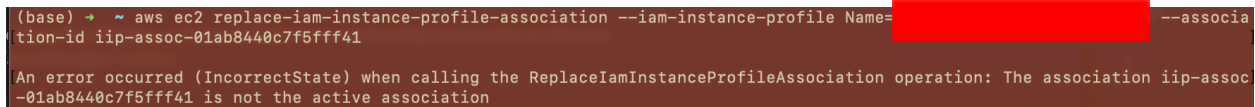
After getting the association ID, attempted the below options to replace but kept running into errors from both methods:

- Replace-iam-instance-profile-association
- Replacing instance profile from AWS console

Error from AWS Console replacement:



Error from CLI replacement:



I ended up having to fully dissociate the role using:
aws ec2 disassociate-iam-instance-profile

Then associated the newly created role with:

Aws ec2 associate-iam-instance-profile

<https://docs.aws.amazon.com/cli/latest/reference/ec2/associate-iam-instance-profile.html>

4.4. Verify Access from the EC2 Instance

Validating access to my bucket:

```
ssh -i /path/to/key.pem ec2-user@<EC2_PUBLIC_IP>
```

After SSHing in, was able to list bucket contents

The image shows a terminal window with the following text: "[ec2-user@ [REDACTED] ~]\$ aws s3 ls s3://[REDACTED]" followed by the output: "2025-06-15 22:50:01 1345438 hisoka_0-0.png".

As I made incremental changes to policy while working through the lab, I used simulate-principal-policy to test access (e.g. in example below, testing ability to get specific object in bucket)

```
aws iam simulate-principal-policy \  
--policy-source-arn arn:aws:iam::476569303606:role/EC2-S3-Access-Role-Brandon \  
--action-names s3:GetObject \  
--resource-arns arn:aws:s3:::discord-labs/hisoka_0-0.png
```

5. Stretch Goals

Create bucket policy that blocks all public access but allows the previously created IAM role

In the lab .md, a policy statement that denies insecure transport (denies all access over HTTP) is provided. I combined the provided policy statement with a statement that allows my IAM role to access the bucket

- Principal: the IAM role
- Action: all s3 commands
 - I later tweaked this to GetObject and ListBucket actions
- Resource: lab bucket
- Condition: requiring secure transport (i.e. over HTTPS), this is redundant when paired with the first policy statement
 - I later removed this due to the redundancy, but unsure if this is best practice

I included Sids for each statement for clarity, and made the changes from the AWS console (bucket -> permissions -> bucket policy -> edit -> pasted in contents from VSCode file shown below)

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "DenyInsecureTransport",
      "Effect": "Deny",
      "Principal": "*",
      "Action": "s3:*",
      "Resource": [
        "arn:aws:s3::[REDACTED]",
        "arn:aws:s3:::[REDACTED]/*"
      ],
      "Condition": {
        "Bool": {
          "aws:SecureTransport": "false"
        }
      }
    },
    {
      "Sid": "AllowAccessToSpecificIAMRole",
      "Effect": "Allow",
      "Principal": {
        "AWS": "arn:aws:iam:[REDACTED]:role/EC2-S3-Access-Role-Brandon"
      },
      "Action": "s3:*",
      "Resource": [
        "arn:aws:s3:::[REDACTED]",
        "arn:aws:s3:::[REDACTED]/*"
      ],
      "Condition": {
        "Bool": {
          "aws:SecureTransport": "true"
        }
      }
    }
  ]
}
```

Restrict to your IP Address:

I created a statement which allows requests coming from my IP to list or read objects in the bucket.

```
{
  "Sid": "AllowGetListFromSpecificIP",
  "Effect": "Allow",
  "Principal": "*",
  "Action": [
    "s3:ListBucket",
    "s3:GetObject"
  ],
  "Resource": [
    "arn:aws:s3:::[REDACTED]",
    "arn:aws:s3:::[REDACTED]"
  ],
  "Condition": {
    "IpAddress": {
      "aws:SourceIp": "[REDACTED]/32"
    }
  }
}
```

I then adapted it for the existing statement (which allowed access to my IAM role), by changing the Principal from * to my IAM role in order to tighten up access.

```
{
  "Sid": "AllowGetListFromSpecificRoleAndIP",
  "Effect": "Allow",
  "Principal": {
    "AWS": "arn:aws:iam::[REDACTED]:role/[REDACTED]"
  },
  "Action": [
    "s3:ListBucket",
    "s3:GetObject"
  ],
  "Resource": [
    "arn:aws:s3:::[REDACTED]",
    "arn:aws:s3:::[REDACTED]/*"
  ],
  "Condition": {
    "IpAddress": {
      "aws:SourceIp": "[REDACTED]/32"
    }
  }
},
{
```

To do the same from CL (taken from lab instructions):

Save the JSON above to a file `bucket-ip-policy.json`, then run:

```
aws s3api put-bucket-policy \ --bucket your-bucket-name \
--policy file://bucket-ip-policy.json
```

Experiment with requiring MFA or VPC conditions

To determine VPC ID from CLI:

```
aws ec2 describe-instances \
  --instance-ids <my instance ID> \
  --query "Reservations[*].Instances[*].VpcId" \
  --region us-east-1
```

```
[
  [
    "vpc-[REDACTED]"
  ]
]
(END)
```

Alternatively, to determine VPC from AWS Console, navigate to the EC2 instance details page:

Instance type
t2.micro

VPC ID
vpc-[REDACTED]

Creating VPC Endpoint (with com.amazonaws.us-east-1.s3 gateway service with existing vpc)

Type [Info](#)
Select a category

☒ **AWS services**
Connect to services provided by Amazon with an Interface endpoint, or a Gateway endpoint

☐ **PrivateLink Ready partner services**
Connect to SaaS services which have AWS Service Ready designation with an Interface endpoint. Uses AWS PrivateLink

☐ **EC2 Instance Connect Endpoint**
An elastic network interface that allows you to connect to resources in a private subnet

☐ **Resources - New**
Connect to resources like Amazon Relational Database Services (RDS) with a Resource endpoint. Uses AWS PrivateLink

☐ **Endpoint services that use NLBs and GWLBs**
Find services shared with you by service name. Connect to a Network LoadBalancer (NLB) service with an Interface endpoint or to a Gateway LoadBalancer (GWLB) service with a Gateway Load Balancer endpoint

Services (1/6)

Q Search

s3 X Clear filters

	Service Name	Owner	Type	Service Region
<input type="radio"/>	com.amazonaws.s3-global.accesspoint	amazon	Interface	us-east-1
<input checked="" type="radio"/>	com.amazonaws.us-east-1.s3	amazon	Gateway	us-east-1

✔ Successfully created VPC endpoint

Notifications 0 0 1 1 0

Endpoints (1) [Info](#)

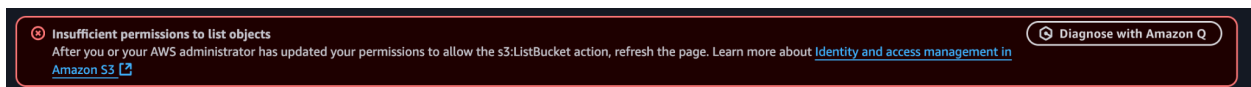
Q Find endpoints by attribute or tag

	Name	VPC endpoint ID	Endpoint type	Status
<input type="checkbox"/>	-		Gateway	✔ Available

After setting up a VPC Endpoint, I decided to edit the bucket policy in a way such that any traffic NOT coming from my VPC was denied, rather than add it as an additional condition to the policy statement permitting access. This will prevent any sources outside of the VPC from accessing the S3 (even if they have the proper IAM role).

```
{
  "Sid": "DenyOutsideMyVPC",
  "Effect": "Deny",
  "Principal": "*",
  "Action": "s3:*",
  "Resource": [
    "arn:aws:s3:::[REDACTED]",
    "arn:aws:s3:::[REDACTED]/*"
  ],
  "Condition": {
    "StringNotEquals": {
      "aws:SourceVpc": "vpc-[REDACTED]"
    }
  }
},
```

An expected side effect of this change is I was no longer able to view the bucket contents from the AWS console:



I also experimented with temporarily changing the condition's SourceVpc value to a different VPC ID to validate that the policy was indeed being evaluated:

```
An error occurred (AccessDenied) when calling the ListObjectsV2 operation: User: arn:aws:sts::476
:assumed-role/EC2-S3-Access-Role-Brandon/i- is not authorized to perform
m: s3:ListBucket on resource: "arn:aws:s3:: with an explicit deny in a resource-bas
ed policy
```

Host a Static Site

TODO: Was not able to get to this, but see below for brief outline and description of steps that I hope to tackle in the future.

Steps

- Create S3 bucket (static file storage)
- Enable static website hosting (serve via HTTP)
- Add public read policy (allow browsers to fetch files)
- Route 53 CNAME or Alias (connect DNS name to S3)
- CloudFront + ACM (Enable HTTPS with SSL)

c. Further EC2 Exploration on Free Tier

1. Snapshots & AMIs

- Create an EBS snapshot of `/dev/xvda`.
 - `/dev/xvda` is the device name assigned to the root EBS (Elastic Block Storage, the storage device) volume of the EC2 instance
 1. Can think of it as main hard drive for EC2 (like C:\ for Windows)
 2. Where OS is installed, typically maps to `/`. usually the first block device attached when launching EC2 instance
 - Then create snapshot from AWS console or CLI
 1. From CLI:
 - a. Get Volume ID (describe-instances with query containing `Ebs.VolumeId`)

```
(base) → ~ aws ec2 describe-instances --instance-id i- --query "Reservations[*]
.Instances[*].BlockDeviceMappings[*].{DeviceName:DeviceName,Ebs:Ebs.volumeId}" --output table
```

DescribeInstances	
DeviceName	Ebs
/dev/xvda	
(END)	

- b. Create snapshot (using create-snapshot of volume ID found above)

```
{
  "Tags": [],
  "SnapshotId": "snap-[REDACTED]",
  "VolumeId": "vol-[REDACTED]",
  "State": "pending",
  "StartTime": "2025-07-06T21:00:14.727000+00:00",
  "Progress": "",
  "OwnerId": "[REDACTED]",
  "Description": "Snapshot of root volume /dev/xdva for discord lab 3",
  "VolumeSize": 8,
  "Encrypted": false
}
```

2. From AWS Console:

vol-

Last updated 2 minutes ago

Actions

Del

Details

Volume ID

vol-

AWS Compute Optimizer finding

Opt-in to AWS Compute Optimizer for recommendations. | [Learn more](#)

Fast snapshot restored

No

Attached resources

i-

dev/xvda (attached)

Source

Size

8 GiB

Volume state

In-use

Availability Zone

us-east-1b

Outposts ARN

-

Type

gp3

IOPS

3000

Created

Wed Jun 11 2025 15:33:15 GMT-0700 (Pacific Daylight Time)

Managed

false

Fault injection

Multi-Attach enabled

No

Operator

-

- Register or create an AMI from that snapshot.
 - Some sources online recommend using create-image command (rather than create-snapshot -> register-image, as snapshots alone don't contain **launch config**, **boot info**, or **kernel info**). But see below for my running of register-image:

```
(base) → ~ aws ec2 create-snapshot --volume-id vol- --description "Snapshot of root volume /dev/xvda for discord lab 3"
```

Register-image (using snapshot ID of snapshot generated above)

```
(base) → ~ aws ec2 register-image \
  --name "Lab3-AMIFromSnapshot" \
  --architecture x86_64 \
  --root-device-name /dev/xvda \
  --block-device-mappings [{"DeviceName": "/dev/xvda", "Ebs": {"SnapshotId": " "}}] \
  --virtualization-type hvm
```

```
{
  "ImageId": "ami-"
}
```


To query for AMI ID (use describe-images to output a table containing image ID):

```
(base) → ~ aws ec2 describe-images \
--owners self \
--query "Images[*].{ID:ImageId,Name:Name}" \
--output table
```

- Lab question: Understand how you can "version" a server with snapshots. Why is this useful?
 - It's useful to 'version' a server using snapshots similar to how it's useful to version your codebase using a code repository. Before carrying out a change, you can take a snapshot of your server for contingency/rollback purposes if something goes astray.
- Launch a new instance from your AMI
 - From CLI:
 1. Utilize run-instances command (passing in the AMI ID into the image-id parameter, security-group-ids, subnet-id, etc)

```
(base) → ~ aws ec2 run-instances \
--image-id [REDACTED] \
--count 1 \
--instance-type t3.micro \
--key-name my-key-pair \
--security-group-ids sg-xxxxxxxxxxxx \
--subnet-id subnet-xxxxxxxxxxxx \
--tag-specifications 'ResourceType=instance,Tags=[{Key=Name,Value=My}'
```

- From AWS Console



2. Linux & Security Tooling

TODO: Did not get to this

- Use `ss -tulpn`, `lsof`, and `auditctl` to inspect services and audit.
- Install and run:
 - `nmap localhost`
 - `tcpdump -c 20 -ni eth0`
 - `lynis audit system`
 - `fail2ban-client status`
 - OSSEC/Wazuh or ClamAV.

3. Scripting & Automation

TODO: Did not get to this

- Bash: report world-writable files (research any commands you don't know!):

```
find / -perm -002 -type f > ww-files.txt
```
- Python with `boto3`: list snapshots, start/stop instances.

8. Reflection

Given my general unfamiliarity with AWS, I had to repeatedly backtrack and review the steps I took earlier in the lab to find incorrect configurations. For example, when testing out policy changes using `simulate-principal-policy` (to test ability to get-object from my s3), I ran into an implicit deny message. I then figured out I incorrectly specified `"arn:aws:s3:::<my bucket name>"` as the resource for the `getobject` action, instead of `"arn:aws:s3:::<my bucket name>/*"`. After making the change, I was able to successfully run `getObject` using `simulate-principal-policy`.

The lab was definitely illuminating in terms of the number of potential security issues that can arise when provisioning/maintaining a larger environment, as security/policy creation took a lot of thought even when working with only 2 components. I faced a few policy design decisions that took me a while to weigh the pros and cons of (e.g. whether to add the IP or VPC condition as part of an 'Allow' action if the source value matched or as part of a 'Deny' if the source value did NOT match). Or as another example, what to define the Principal value as for a statement (either as `*` or as the role assigned to the EC2 to balance security vs availability).

As part of the lab, a key realization I made and (hope to drill into my brain haha) is that each statement in a bucket policy is evaluated independently, and that permissions are cumulative (unless there's an explicit Deny).